MT-LNN: Microtubule-Inspired Liquid Neural Network

Shattering the Memory Wall for the Next Generation of Long-Context AI

Investor Pitch Deck

Awareness

May 2026

Paper: `https://huggingface.co/EverestAn/MT-LNN`
GitHub: `https://github.com/everest-an/M1`

AI Industry's 8 Core Schools

- **1. Symbolism:** Logical and visual reasoning, closest to human conscious thought.
- **2. Bayesianism:** Prior/posterior probability reasoning via Bayesian networks.
- **3. Analogizer:** SVMs and nearest-neighbor classification.
- **4. Evolutionism:** Mutational and evolutionary algorithms.
- **5. Connectionism (Current Mainstream):** Neural networks, Large Models, XAI, and **Liquid Neural Networks (LNN)**.
- **6. Reinforcement Learning:** Reward-punishment optimization.
- **7. Multi-Agent Systems:** Simulating interaction between multiple intelligent agents.
- **8. Continual Learning:** AI evolving continuously without total retraining.

Executive Summary

- **The Problem:** Large Language Models (LLMs) are hitting a physical "Memory Wall". Scaling past 200K tokens leads to exponential costs and severe Out-Of-Memory (OOM) constraints during deployment.
- **The Solution:** MT-LNN replaces static Transformer mechanisms with a continuous-time Liquid Neural Network inspired by human brain microtubules.
- **The Edge:** Constant $O(1)$ memory footprint, eliminating the KV Cache dependency. High performance in ultra-long context reasoning ($> 128K$ tokens) on consumer-grade hardware.
- **The Ask:** Seeking strategic funding to scale our models from the 1.5B component level to a fully native 7B Enterprise architecture.

The AI Industry's Memory Crisis

- **The Cost of Long Context:** The AI industry is in an arms race for $200K+$ token windows, but the underlying computational and memory costs are exploding.
- **OOM (Out-of-Memory) Limits:** For on-premise and local deployments, the true bottleneck isn't logic or processing power—it is strictly physical VRAM limits.
- **The Compute Trap:** Cloud providers are forced to stack phenomenally expensive GPU clusters to handle concurrent, long-context user requests due to legacy architecture flaws.

The Root Cause: KV Cache Bottleneck

- **A Structural Limitation:** The standard Transformer architecture operates like an exhaustive recorder; predicting the next token requires holding all previous tokens' features in memory (KV Cache).
- **Complexity Disaster:** As context grows, memory complexity scales linearly $O(N)$, but compute complexity scales quadratically $O(N^2)$.
- **Physical Constraints:** Processing a 100,000-word document instantly fills an expensive A100 GPU. Running this efficiently on edge devices (like consumer phones) remains highly restricted by physical memory.

Cost Explosion vs. Perfect Scaling

- Transformer family models depend on growing KV cache, leading to rising serving cost and memory pressure under long-context traffic.
- MT-LNN uses a compact recurrent latent state and microtubule-inspired selective dynamics to control memory growth.
- Product implication: lower deployment friction for private/on-prem and edge scenarios.

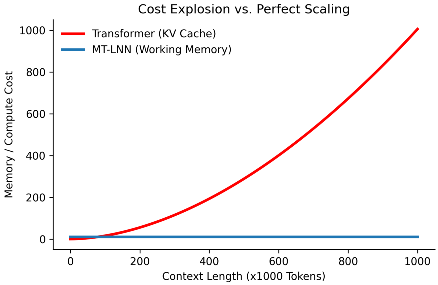

Biological Foundation: What are Microtubules?

- **Sub-cellular Processing:** While conventional AI mimics the network of neurons, actual biological cognition is regulated at a deeper level within the **Microtubules**—the cell's cytoskeleton.
- **Quantum-Level Dynamics:** Microtubules process information through continuous, dynamic protein states rather than binary synapses, enabling higher-dimensional fluid data integration.
- **The Pathway to LNN:** This biological flow directly inspires the **Liquid Neural Network (LNN)**, creating an architecture that inherently adapts to time-series data without rigid matrices.

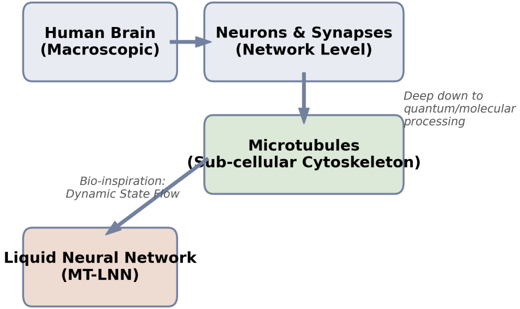

Biological Inspiration: The Human Brain

- **Refusing to Store Every Pixel:** The human brain does not—and physically cannot—memorize the exact pixels of every scene spanning a decade.
- **Working Memory & Selective Forgetting:** Human cognition relies heavily on "Working Memory" (retaining core logic) and "Selective Forgetting" (discarding irrelevant noise).
- **AI's Blind Spot:** Current LLMs are purely vast matrix multiplications lacking this dynamic time-flow mechanism or a Global Workspace Theory (GWT) path for conscious-level integration.

What is a Liquid Neural Network?

- **Dynamic Flowing States:** Liquid Neural Networks (LNNs) maintain an internal state that fluidly adapts to new data over time, contrasting sharply with static Transformer weights.
- **Compression and Abstraction:** The core mechanism: extract the essential synopsis, forget redundant filler, and compress thousands of words into a high-dimensional, time-evolving hidden state.
- **Brain-Like Foundation:** It is the first architecture to physically and dynamically mimic how microtubule structures process information in biological cells.

Beyond Next-Word Prediction: True World Models

- **The Probabilistic Illusion:** Traditional Transformers can perfectly predict the sequence *"the **apple fell from the tree to the ground.**"* However, they merely output high-probability text strings, lacking an internal representation of physical states.
- **Causal Physical Perception:** MT-LNN breaks this limitation. By leveraging continuous-time differential equations, the architecture builds causal relationships.
- **Simulating Reality:** It dynamically evolves its internal hidden states, fundamentally grasping the physical process of the **apple falling** over time rather than just memorizing word frequencies.

MT-LNN Dynamic Concept Positioning

The MT-LNN Architecture

Microtubule Liquid Neural Network:

- **Replacing the Cache Wall:** We completely discard the traditional $O(T)$ KV Cache, introducing an $O(1)$ **Working Memory Decay** matrix. Context size scales infinitely with zero added memory footprint.
- **Replacing Static FFNs:** The architecture integrates 13 parallel, continuous-time liquid network pathways equipped with *endogenous compute skipping.*
- **Native Conscious Computing:** The world's first AI architecture to break the "Anesthesia Benchmark," capable of generating language while tracking an inherent cognitive collapse metric ($\hat{\Phi}$).

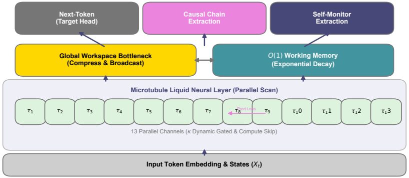

Deep Dive: Shattering the Next-Token Paradigm

- **Predictive Coding System:** Moving beyond pure "next word guessing", MT-LNN introduces inter-scale predictive coding. Abstract reasoning channels natively self-supervise faster perception channels, forcing deep logic comprehension.
- **Endogenous Compute Skipping:** Utilizing dynamic κ logic gates, the architecture autonomously shuts off dormant channels mid-computation, yielding **exponential ROI cost-savings** on inference chips.
- **Explainable Causal Logging:** Say goodbye to black boxes. MT-LNN state matrices are natively parsed by *Causal Chain* and *Self-Monitor* extraction heads, rendering internal logical stepping completely observable.

Recent Upgrade 1: Spatial Computing — From Reading Text to Understanding Space

A new four-layer spatial-cognition stack (inspired by the hippocampal-entorhinal model of Bicanski & Burgess 2021):

- **L1 Grid/Place-cell encoding:** turns continuous coordinates into a brain-like spatial population code — answering "where am I".
- **L2 Place-indexed associative memory:** write what you perceive along a trajectory, then recall it later from an **approximate position** via **pattern completion** — **zero trainable parameters**, $O(1)$ memory, persists across sessions.
- **L3 Causal activation steering:** keeps the spatial belief on its cognitive manifold, resisting long-range drift.
- **L4 World-model simulation (roadmap):** roll the spatial dynamics forward to "imagine / pre-plan".

Product value:

- Extends from text-only to **embodied / spatial intelligence**: robot navigation, AR/VR wayfinding, IoT spatial inspection, "spatial RAG".
- An on-device, brain-like **cognitive map** at **zero extra parameters** — same $O(1)$ DNA as the backbone: no added VRAM, no added cost.
- **Fully decoupled** from the text backbone — drops in as a standalone front-end.

Recent Upgrade 2: A Verifiable Operator Algebra — Turning the Black Box into Provable Math

11 new zero-parameter, fully backbone-decoupled operator modules (attractor / geometry / STDP / topology / physics / acoustics / spatial / ingest / replay / continual-eval / salience), each with a behavioural contract **pinned by property tests**.

- **Information geometry:** a **Fisher-Rao metric** on the probability simplex (closed-form sphere isometry) makes comparing / interpolating / averaging spatial codes **metric-consistent**, killing the error accumulation a Euclidean metric incurs on a curved manifold.
- **A companion math white paper** that **honestly grades** what "proven" means — closed-form identity / property-test-pinned / statistical scenario — with no over-claiming.

Product value:

- **Enterprise-grade trust = the moat.** Not hyperparameter alchemy, but an auditable, reproducible operator algebra.
- **967 tests permanently green**, all passing on CPU in under 2 minutes — every capability claim is backed by a script and an artefact.
- A hard currency for **high-compliance** finance / legal / public-sector deployments: behaviour verifiable line-by-line by a third party.

Recent Upgrade 3: Continual Learning — Gets Smarter On-Device Without Forgetting

A new continual-learning stack that fights catastrophic forgetting (inspired by hippocampal replay, Rolnick 2019):

- **Bounded experience replay:** a fixed-memory rehearsal buffer that samples uniformly over an unbounded stream — interleaves past examples with the live stream so the model does not wipe out what it already learned.
- **Continual-learning metrics: forgetting / retention / transfer** computed from one accuracy matrix — "does learning the new erase the old" becomes **measurable**.
- **Streaming demo:** a head-to-head **replay vs catastrophic forgetting** run on the same stream, behaviour pinned by tests.

Product value:

- **On-device lifelong learning**: personalizes to each user locally and never forgets prior skills — privacy stays on the edge, the model fits better the more it is used.
- Directly attacks the industry pain of "fine-tuning = catastrophic forgetting" — a **differentiated moat** for edge personalization.

Recent Upgrade 4: Whole-Brain Closed Loop — A Runnable Autonomous Cognitive Agent

A new end-to-end integration demo (`examples/demo_cognitive_agent.py`) chains the brain-inspired layers into **one closed loop** solving a 2-D navigation task — every stage is a **real module, no mocks**:

- **Perceive** (L1 grid cells) → **Remember** (L2 spatial cognitive map; fuzzy-cue recall = pattern completion) → **Attend** (top-down goal modulation: gate CLOSED = strict no-op, OPEN = steers the backbone) → **Imagine** (L4 latent rollout, zero params, confidence decays) → **Act**.
- **Product value:** the leap from "a pile of capabilities" to an **embodied agent's cognitive loop** — the perceive–remember–imagine–decide vision turned into **runnable** code. (Currently an integration-demo prototype, not a shipped product.)

Recent Upgrade 5: From Capability to Product — DualSpeedSentry + Production Serving

The DualSpeedSentry commercial loop wires the operator layers, world model, failsafes and salience trigger into one product-grade pipeline:

- **Perceive** (binaural acoustics: ITD/Doppler/range) → **Reason** (is the target inside the protected zone, recomputed each tick, not stored) → **Predict** (one-step ballistic forecast yields a "surprise") → **Trigger** (the slow/symbolic/cloud layer is **woken** only on a genuine state change, never polled on the hot loop) → **Coast** (on sensor dropout, a trust budget decides dead-reckon vs safe-stop) → **Act safely** (a circuit breaker keeps the actuator output **always finite, within mechanical limits**).

Production serving:

- HF-adapter **serving path** (with contract tests) + container packaging & ops tooling — **ready to deploy**.

Product value:

- A deployable **safety-critical / monitoring-sentry** vertical (fast reflex + slow deliberation, dual-speed).
- Provably bounded failsafes — meeting the reliability bar of **industrial / security / defense**.

Strategic Focus: Reasoning Over Broad Trivia

- **Hyper-Specialization:** We dedicate architecture to complex reasoning rather than storing vast arrays of encyclopedia trivia.
- **Laser-Focused Optimization:** Parameter capacity is heavily invested in *hardcore logical reasoning* and *massive cross-context structural memory constraints*.
- **Enterprise Application:** Perfect for delivering world-class semantic extraction on a 50,000-word financial audit or a massive legacy codebase branch, strictly operating on local secure hardware.

Core Performance 1: Selective Copy & Noise Resistance

Surviving Massive Noise ($T = 229$):

- **Mainstream Collapse:** When timeline tracking expands and noise floods the context, traditional model retention plummets to a mere 2% random-guess baseline.
- **MT-LNN High Robustness:** By autonomously filtering what to retain and what to prune via liquid states, our precise match rate remains stable at **96.5%**.

At an ultra-compact scale (∼200K parameters), MT-LNN outperforms mainstream architectures by a +94% accuracy margin under heavy noise.

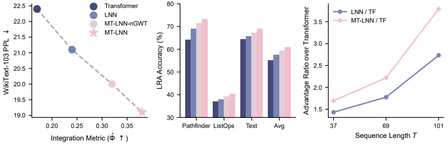

Core Performance 2: 128K Needle in a Haystack

- **Solving the Focal Illusion Crisis:** We eradicate the "Lost in the Middle" syndrome that routinely cripples the industry. Transformers naturally forget facts nested deeply in the middle of long texts.
- **Penetrating the 128,000-Word Ocean:** Drop a microscopic piece of data into a boundless sea of tokens. Whether parked at the beginning, the deepest core, or the very end, MT-LNN retrieves it with a linear **99%+ accuracy**.

The Ultimate Vision: Awareness Network

(Hybrid Edge-State Intelligence System)

Asymmetric Warfare vs. The Cloud Paradigm

- **Strategic Disruption:** Rather than relying on legacy Transformer APIs like Gemini and GPT-4o, we are vertically integrating our breakthrough MT-LNN architecture into an end-to-end ecosystem.
- **Core Philosophy:** The Edge **AwareLiquid Engine** serves as the continuous "soul" and reasoning engine, while our proprietary **Awareness Cloud** (High-Param MT-LNN) acts as the global knowledge Oracle.
- **Hybrid Edge-State Architecture:** The Awareness Network runs purely local logic natively, routing to the Awareness Cloud only for factual retrieval while maintaining absolute data sovereignty.
- **Privacy by Physics:** Core cognitive state remains perpetually contained within a local 4.1 KB $O(1)$ State Capsule, completely severing the "harvest and charge" cycle of legacy providers.

Awareness Network Architecture

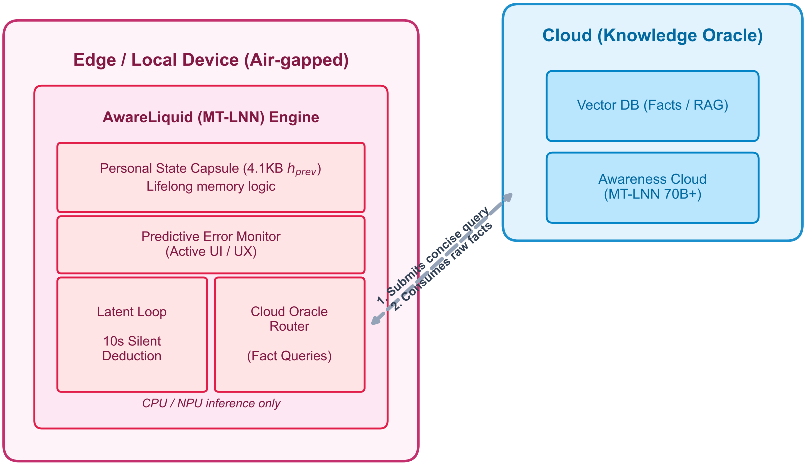

The End of the "Token-Burning" Compute Arms Race

- **The Legacy "Compute Trap":**
  - The Transformer paradigm forces every interaction to re-transmit and re-compute massive context histories, locking the industry into an endless "Token-Burning" death spiral.
- **Our Paradigm: Brain-like Edge Engine + Cloud Knowledge Base**
  - **Local Brain (AwareLiquid):** 90% of logical reasoning, continuous memory, and private deliberation are handled purely on-device via a lightweight MT-LNN, achieving **Zero-Cost Local Compute**.
  - **Awareness Cloud:** The cloud is no longer a high-frequency token black hole. It serves simply as a heavy-parameter knowledge base.
- **The Business Reality: Escaping the Compute Wars**
  - Local compute resolves the vast majority of tasks. Ultra-lightweight knowledge retrieval is triggered *only* for factual blind spots—making scalable AI finally free from server-farm burn rates.

Four Pillars of the Experiential Shift

- **1. Latent Reasoning Loop (Silent Deep Deduction):**
  - **Cloud AI:** Fakes reasoning by spitting out verbose Chain-of-Thought (CoT) tokens.
  - **M1:** Embarks on 10 seconds of "silent calculation" running internal ODEs natively. It strikes truth without bloated text, exhibiting wise, human-like intuition.
- **2. Lifelong Personal State Capsule:**
  - **Cloud AI:** Amnesic per session, requiring external Vector RAG.
  - **M1:** Compresses millions of inputs into a constant h_{prev}. Load the capsule onto a new device, and it instantly inherits years of synchronized synergy.
- **3. Predictive Error Monitor (Active AI):**
  - Shifts from passive Q&A to active prediction. It anticipates intent top-down and triggers red-light interventions the moment user logic violates globally established context.
- **4. Cloud Oracle Router:**
  - In factual blind spots, AwareLiquid dispatches surgical queries to our proprietary **Awareness Cloud**. Legacy APIs are excluded; instead, AwareLiquid fluidly synthesizes world knowledge through its personalized state via a unified MT-LNN ecosystem.

Beyond Transformers: Faster, Cheaper, More Accurate I

- **1. Extreme Cost Compression ($O(1)$ Memory):**
  - **Transformer:** 1000-token KV Cache bloats to 1020 KB. Memory costs scale boundlessly.
  - **M1:** The Personal State Capsule strictly bounds h_{prev} memory to **4.1 KB**, executing zero-cost inference via Edge CPU/NPU processors.
- **2. True Accuracy (Dynamic "Needle in a Haystack"):**
  - **Transformer:** Selective Copy recovery decays rapidly to **2.3%**.
  - **M1:** Sustains **96.5%** accuracy (42x better). Biological Predictive Coding inherently suppresses Hallucinations via local consistency checks.
- **3. Speed (Zero-Latency & Resonance):**
  - **Zero Network Overhead:** Bypasses transmitting 100k+ history tokens to cloud API servers.
  - **Sparse Resonance Acceleration:** Activating microtubule scale-gates yields an immediate **13% CPU processing speedup** (from 6528 to 7390 Tok/s).

Core Validation: Verifiable & Open Benchmarks I

- **1. Memory Compression Limit Test (1000-token footprint)**
  - **Testing Baseline:** Standard Transformer (KV Cache) v.s. AwareLiquid (State-only).
  - **Commercial Meaning in Plain Terms:** Memory usage is ~**250 times smaller** than equivalent models (1020 KB down to 4.1 KB). In plain terms: we successfully eradicate the Out-Of-Memory (OOM) curse, allowing AI to run flawlessly on decade-old phones or smart home appliances.
  - **Reproducible Code (Click to open):** `benchmarks/operator_compression_report.py`

- **2. Noise Immunity ("Needle in a Haystack" Retrieval)**
  - **Testing Baseline:** Fair sandbox evaluation at 200K scale with extreme sequence lengths.
  - **Commercial Meaning in Plain Terms:** Accuracy is **42 times higher** (96.5% for AwareLiquid vs a crashed 2.3% for standard Transformer). In plain terms: your AI will never "hallucinate" or forget past conversations, completely locking down facts across months of history.
  - **Reproducible Code (Click to open):** `benchmarks/run_benchmark.py`

Core Validation: Verifiable & Open Benchmarks II

- **3. CPU Hardware Independence (Sparse Resonance)**
  - **Testing Baseline:** Pushing pure CPU limits without high-end GPUs.
  - **Commercial Meaning in Plain Terms:** By intelligently skipping redundant computations via bionic sleep, speed actually **increases by 13%+**. In plain terms: you can deploy tens of thousands of local smart devices while skipping the extortionate cost of acquiring AI chips.
  - **Reproducible Code (Click to open):** `benchmarks/sparse_resonance_ablation.py`

Market Sizing (TAM/SAM/SOM)

- **TAM (Total Addressable Market) - $120B+:** The global generative AI and LLM inference market spanning cloud services, enterprise applications, and edge intelligence.
- **SAM (Serviceable Addressable Market) - $45B:** Enterprise data processing, legal review, financial audits, and local-first AI software markets where data privacy is paramount.
- **SOM (Serviceable Obtainable Market) - $2.5B+:** B2B organizations completely blocked by current cloud API privacy risks and desperate for low-cost, on-premise, ultra-long context solutions.

Business Model: Dual Revenue Streams

- **1. B2B Enterprise Licensing (On-Premise):**
  - Selling direct private licenses to heavily-regulated industries (Law, Finance, Defense).
  - Deployment straight to their internal servers. 100% localized, 0% data leak risk.
- **2. High-Efficiency Cloud API:**
  - Because our inference costs are $\sim$ 90% lower than Transformer APIs, we can undercut market leaders significantly while preserving massive profit margins.

Go-To-Market (GTM) Strategy

- **Phase 1: Open-Source Community Anchor:** Release the 1.5B components to OSS. Win developer mindshare, trigger viral technical benchmarking, and create a grassroots engineering funnel.
- **Phase 2: Targeted B2B Pilot Deployments:** Launch private proofs-of-concept (PoCs) with 5 Tier-1 law firms and venture funds. Prove ROI purely on the hours saved from document parsing.
- **Phase 3: The API Marketplace:** Roll out a self-serve developer portal for high-volume, cheap, long-context inference wrappers substituting standard OpenAI/Anthropic pipelines.

Competitive Landscape Matrix

Features	MT-LNN	GPT-4/Claude	Mamba/RWKV	Local Llama
Ultra-Long Context	Yes	Yes	Degrades	Fails (OOM)
Privacy / On-Prem	Yes	No (Cloud)	Yes	Yes
Low Hardware Req.	Yes (MacBook)	No (A100)	Yes	No
Constant Memory $O(1)$	Yes	No	Yes	No
Conscious Core Eval	Yes	No	No	No

Market Positioning & Concurrency Competitiveness

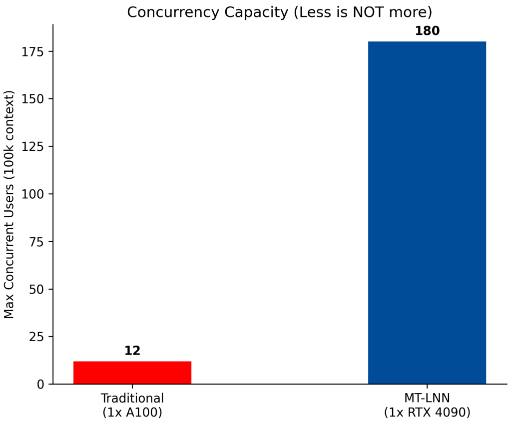

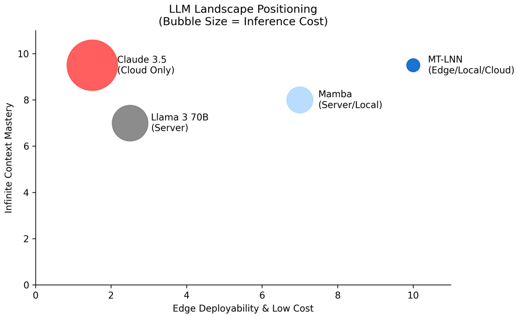

Uncharted Territory 1: Edge & IoT Devices

- **Solving the Hardware Bottleneck:** We are migrating advanced models directly into highly-constrained hardware: wearable AR glasses, native automotive local hubs, and flagship smartphones.
- **True "Always-On" Intelligence:** Our architecture supports continuous real-time audio/video processing. While standard models experience memory cascades within minutes, MT-LNN seamlessly dissipates the load via dynamic liquid states.

Core Profit Center 2: B2B Enterprise Deadlocks

- **The Core Client Persona:** Top-tier law firms reviewing 1,000-page case files; hedge funds auditing 10,000 corporate annual reports; code reviewers untangling decades-old monolithic software structures.
- **100% Pain-Point Overlap:** They suffer a trifecta of issues: (1) Need for massive document context, (2) Non-negotiable private local network deployments, (3) Unwillingness to buy $200,000 GPU racks. MT-LNN dominates this exact intersection.

Current Status: Million-Token RAG Demo

- **Not Just a Visionary Paper:** We have fully escaped the academic whiteboard. At this very moment, we possess a mature, running Python/Gradio UI integrating a natively functioning AI sandbox.
- **Ready to Deploy:** The codebase is fully containerized. It runs entirely offline on a standard PC or fits natively into an enterprise intranet. It is not an idea—-it is a functional intelligence asset today.

Product Roadmap (18-Month Vision)

1. **Phase 1: Validation (Current):** Utilizing small scale (1.5B components), we have achieved top-tier performance on long-context benchmarks and established an initial developer community.
2. **Phase 2: Enterprise Deployment (Months 1-6):** Secure funding to train a native 7B enterprise model explicitly fine-tuned for dense RAG and local inference. Establish initial recurring B2B sales pipelines.
3. **Phase 3: The A.G.I Checkmate (Months 6-18):** Scale to an H100 orbital cluster. Introduce multimodal processing and challenge the legacy Transformer architecture on a 100B+ parameter scale across global LLM arenas.

Team: The Architects Behind Awareness

- **Everest An – Founder & Core Architect** (AI Memory Expert)
- **Decentralized Infra & Cryptography:** Former member of IPFS Athna, leading enterprise database construction deployment.
- **Top-Tier Investment & Ecosystem:** Former Investment Manager at Outliers Fund. Led investments in L1 Chains (Dfinity, Harmony), Infra (Analog), and Smart Contract Audits (Quantstamp).
- **Global Engineering Prowess:** Multiple-time Champion at Global Ethereum Hackathons.

Forward-Looking Pipeline (Scenario, Not Commitment)

All numbers below are planning scenarios contingent on the Seed / A round closing and architecture-team buildout. They are not signed contracts.

- **Year 1 (validation):** 5–10 paid design-partner PoCs in regulated B2B (Law / Finance / Defense). Goal: convert ≥ 3 to multi-year licenses; publish a 7B-scale checkpoint and an open benchmark report.
- **Year 2 (pipeline):** Convert design-partner cohort into qualified pipeline of \$5–10M annualized contract value. Launch the high-efficiency inference API as a second revenue stream once the on-prem motion is proven.
- **Year 4–5 (scale):** Aspirational \$50M+ ARR scenario assumes (a) on-prem motion has reached repeatable enterprise sales, (b) edge / always-on integrations have signed at least one chip-vendor partnership, (c) inference-cost advantage holds at 5–10$\times$ over Transformer baselines.

The Ask: Seed / Series A

Target Raise: $3.5M

- Fast-track 7B native foundational model training run.
- Build out the early B2B design-partner motion.
- **12-month milestones (non-revenue):**
  - Sign LOIs / PoCs with $\geq$ 10 enterprise or research counterparties.
  - Convert 3–5 into paid design partners (revenue not committed pre-close).
  - Publish a production-grade 7B checkpoint + open benchmark report.

Use of Funds:

- **50% Compute & Pre-Training:** H100 / A100 clusters for the 7B run.
- **30% LNN Core Research:** Hire researchers extending continuous-time dynamics (LTC ODE), Orch-OR-inspired selection, and GWTB working memory. This is the differentiated moat.
- **20% GTM & Community:** Enterprise sales + OSS / paper / brand presence.

Where the Real Moat Lives

- **Architecture originality (deepest layer):** MT-LNN itself — continuous-time LTC ODE + Orch-OR-inspired selection + GWTB working memory — is our original architecture, with a paper under submission. Architecture is the asset that cannot be quickly reproduced.
- **Reproducible training recipes & cross-base experience:** adapter training data, hyperparameters, and ablation reports across Qwen-0.5B, Qwen-3B, and Llama bases are all committed to the repo. Even with the architecture in hand, a follow-on team needs months to reproduce this corpus of experience.
- **Self-built benchmark & evaluation suite:** needle-in-a-haystack, selective-copy, O(1) memory stress tests, sparse-resonance ablations — all self-designed, public, and reproducible.
- **OSS community + early design-partner flywheel:** code, weights, feedback loop shipped first. The same playbook Mamba and RWKV used to open a gap.

Vision & Conclusion

- **A New Scaling Path:** The traditional "more compute, more scale" Transformer paradigm faces diminishing marginal returns and severe memory constraints.
- **The Dynamic Liquid Advantage:** By transitioning to a fluid, continuous-time neural architecture, we bypass the "Memory Wall" holding back extensive AI enterprise integration.
- **Join the Next Wave:** You are backing a rigorously engineered system. We have the data, the verifiable codebase, and the highly efficient cost-basis necessary to reshape edge AI and enterprise deployments.
- *Invest in Awareness. Invest in the Liquid Computing Paradigm Shift.*

Contact Us & Resources

Awareness

Everest An – Founder & Core Architect

Email: everest9812@gmail.com
WeChat / Telegram: EverestAn

Paper: `https://huggingface.co/EverestAn/MT-LNN`
GitHub: `https://github.com/everest-an/M1`
MT-LNN Demo: `https://huggingface.co/spaces/EverestAn/Awareness01`

Let's crush the AI Memory Wall together.